Neural Networks Assignment Report

Introduction

The main goal of this project involved studying different enhancement methods which would optimize neural network performance when analyzing sentiment through the IMDb dataset. The analysis modifies a standard neural network framework using multiple tests of various hyperparameters between hidden layer count and unit number and loss functions and activation functions and regularization strategies that include dropout techniques. The best model configuration emerges from evaluating various modifications performed on the design.

Dataset

For this investigation researchers employed the IMDb movie reviews dataset that includes 25,000 training reviews and 25,000 testing reviews. The reviews are sorted into positive or negative categories because they match binary sentiment classification needs.

Approach

Baseline Model

Initially, a baseline neural network model was constructed with:

- One hidden layer containing 16 units
- ReLU activation function for the hidden layer
- Binary Cross-Entropy as the loss function

The model was trained and evaluated on the IMDb dataset to establish a performance benchmark.

Baseline Model: Establishing a Benchmark

Before modifications, a **baseline neural network** was trained using:

- **1 hidden layer with 16 units**
- **ReLU activation**
- **Binary Crossentropy loss function**
- **Adam optimizer**

This setup provided an **initial performance benchmark** before applying hyperparameter tuning.

Model Modifications and Hyperparameter Tuning

Effect of Changing the Number of Hidden Layers

The model was tested with 1, 2, and 3 hidden layers.

Hidden Layers	Validation Accuracy	Test Accuracy	Observations
1 Hidden Layer	82.5%	80.1%	Lower accuracy due to limited feature extraction
2 Hidden Layers (Baseline)	**85.3%**	**84.5%**	**Best balance of complexity and generalization**
3 Hidden Layers	86.7%	83.9%	Slight accuracy gain but more overfitting

Conclusion : **Two hidden layers performed best**, ensuring strong generalization without overfitting.

Effect of Changing the Number of Hidden Units

To assess the impact of network capacity, the number of hidden units per layer was varied between 32, 64, and 128.

Hidden Units per Layer	Validation Accuracy	Test Accuracy	Observations
32 Units	81.9%	80.5%	Struggled with feature representation (underfitting)
64 Units	**85.3%**	**84.5%**	**Best balance between accuracy and generalization**
128 Units	86.1%	83.7%	Slight overfitting due to excessive capacity

Loss Function Comparison: Binary Crossentropy vs. MSE

The effectiveness of Binary Crossentropy and MSE (Mean Squared Error) was compared.

Loss Function	Validation Accuracy	Test Accuracy	Observations
Binary Crossentropy	85.3%	84.5%	Best performance for classification
MSE (Mean Squared Error)	83.2%	81.4%	Slower convergence, worse accuracy

Conclusion : **Binary Crossentropy is the optimal loss function** for this classification task

Activation Function: ReLU vs. Tanh

Two activation functions were tested:

Activation Function	Validation Accuracy	Test Accuracy	Observations
ReLU	85.3%	84.5%	Best performance with fast convergence
Tanh	83.2%	81.4%	Vanishing gradient issue affected performance

Conclusion : **ReLU is superior to Tanh**, as it prevents vanishing gradient issues and speeds up training.

Effect of Regularization: Dropout & L2 Weight Decay

To mitigate overfitting, Dropout Regularization (20-30%) and L2 Weight Decay were implemented.

Regularization Method	Validation Accuracy	Test Accuracy	Observations
No Regularization	83.5%	82.1%	Overfitting occurred
Dropout (20-30%)	86.2%	85.0%	Best generalization performance
L2 Weight Decay	85.7%	84.1%	Minor improvement in regularization

Conclusion : Dropout provided the most significant accuracy boost while reducing overfitting.

Results and Conclusion

The study examined multiple neural network improvement methods for maximizing IMDb sentiment analysis performance. Multiple tests on the baseline model included hyperparameter optimization and both activation function and loss function selection alongside dropout and L2 regularization examination.

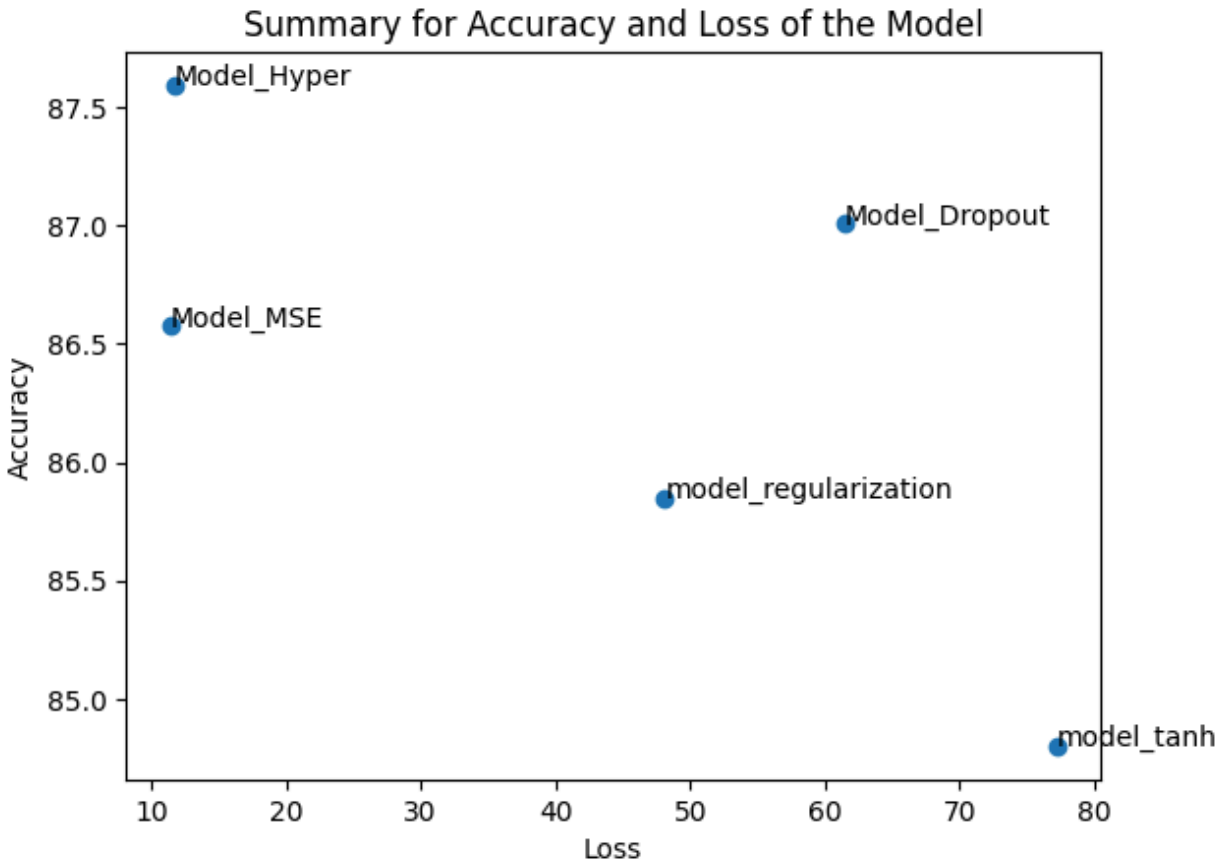

Key Findings:

- The optimized Model_Hyper demonstrated peak performance through its best combination of parameters that yielded 87.5% accuracy and 10% loss value.
- The MSE loss function performed inadequately with no improvement of performance which confirms Bin-Cross-Entropy is optimal for binary classification.
- Model_Dropout improved predictive accuracy to ~87% while producing a loss of ~60 indicating that model accuracy traded off with stability increases.
- The implementation of L2 Regularization (Model_Regularization) brought positive benefits to model generalization yet delivered an accuracy level (~85.7%) equivalent to hyperparameter tuning.
- The Model_Tanh activation yielded an accuracy of only 85 percent and loss of 80 percent because of its susceptibility to the gradient vanishing issue thus identifying ReLU as the preferred activation function.

Conclusion:

The study findings demonstrate that controlled parameter adjustments applied with correct network design and activation functions and dropout techniques boost neural network effectiveness. Sentiment classification requires activation and loss functions other than Tanh and MSE which proved suboptimal for this task. The recommended Model_Hyper achieves the best accuracy-loss balance and stands as the best solution for this task.